

Algorithms and Datastructures

Cache Efficiency, Divide and Conquer

Albert-Ludwigs-Universität Freiburg



**UNI
FREIBURG**

Prof. Dr. Rolf Backofen

Bioinformatics Group / Department of Computer Science
Algorithms and Datastructures, March 2018

Cache Efficiency

- Introduction

- Cache Organization

Divide and Conquer

- Introduction



Background:



Background:

- Up to now we always counted the **number of operations**
- Assuming this is a good measure for the runtime of a algorithm/tool



Background:

- Up to now we always counted the **number of operations**
- Assuming this is a good measure for the runtime of a algorithm/tool
- Today we will see examples where this is not suitable

Example:

- We sum up all elements of a field a of size n in ...
 - natural order:

$$\text{sum}(a) = a[1] + a[2] + \dots + a[n]$$

- random order:

$$\text{sum}(a) = a[21] + a[5] + \dots + a[8]$$

Python:

```
def init(size):  
    """Creates the dataset."""  
  
    # use system time as seed  
    random.seed(None)  
  
    # set linear order as accessor  
    order = [a for a in range(0, size)]  
  
    # init array with random data  
    data = [random.random() for a in order]  
  
    return (order, data)
```

Python:

```
def run(param):  
    """Processes the dataset."""  
  
    # unpack data  
    (order, data) = param  
  
    # init the sum value  
    s = 0  
  
    for index in order:  
        s += data[index]  
  
    return s
```




Figure: Summing elements in linear order

```
def init(size):  
    """Creates a randomly ordered dataset."""  
  
    # use system time as seed  
    random.seed(None)  
  
    # set random order as accessor  
    order = [a for a in range(0, size)]  
    random.shuffle(order)  
  
    # init array with random data  
    data = [random.random() for a in order]  
  
    return (order, data)
```



Figure: Summing elements in random order

Conclusion:

Conclusion:

- The number of operations are identical for both algorithms

Conclusion:

- The number of operations are identical for both algorithms
- Accessing elements in random order takes a lot longer (Factor 10) **Why?**
- The costs in terms of memory access are very different

Cache Efficiency

Introduction

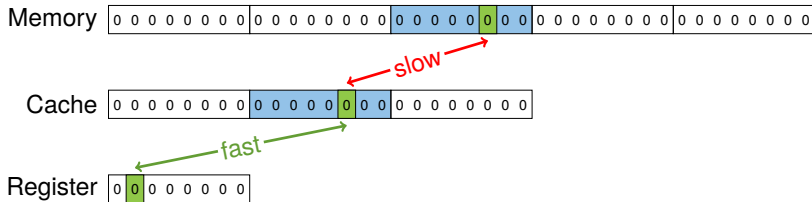
Cache Organization

Divide and Conquer

Introduction

Cache Efficiency

CPU Cache

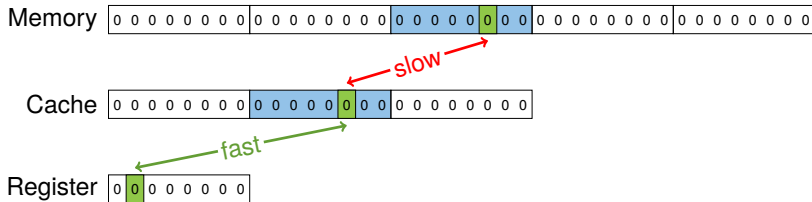


Cache Efficiency

CPU Cache



Principle / organization:



Principle / organization:

- Accessing one byte of the main memory takes ≈ 100 ns



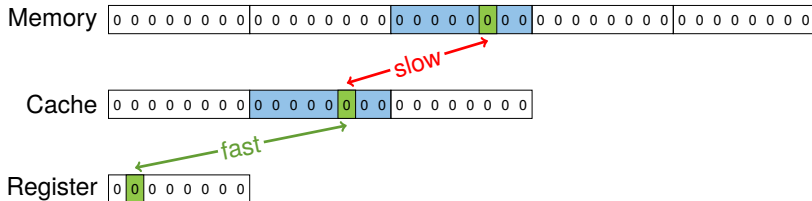
Principle / organization:

- Accessing one byte of the main memory takes ≈ 100 ns
- Accessing one byte of (L1-)cache takes ≈ 1 ns



Principle / organization:

- Accessing one byte of the main memory takes ≈ 100 ns
- Accessing one byte of (L1-)cache takes ≈ 1 ns
- Accessing one or more byte/s of main memory loads a whole block ≈ 100 B into the cache

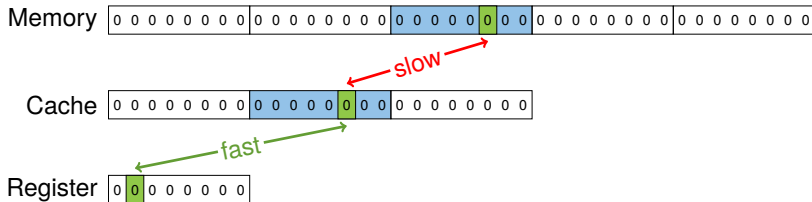


Principle / organization:

- Accessing one byte of the main memory takes ≈ 100 ns
- Accessing one byte of (L1-)cache takes ≈ 1 ns
- Accessing one or more byte/s of main memory loads a whole block ≈ 100 B into the cache
- As long as this block is in the cache, it is not necessary to access the memory for bytes of this block

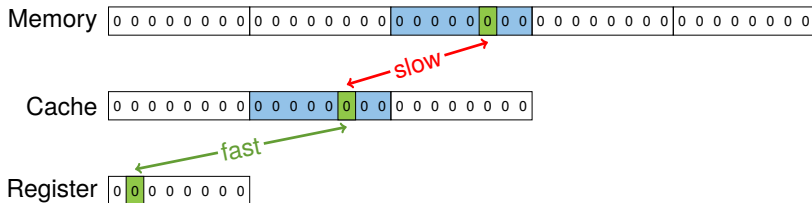
Cache Efficiency

CPU Cache



Cache Efficiency

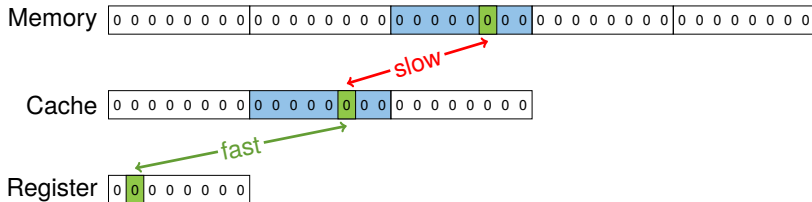
CPU Cache



Cache organization:

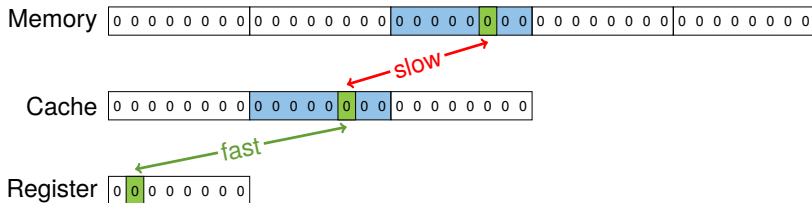


- 13 / 51



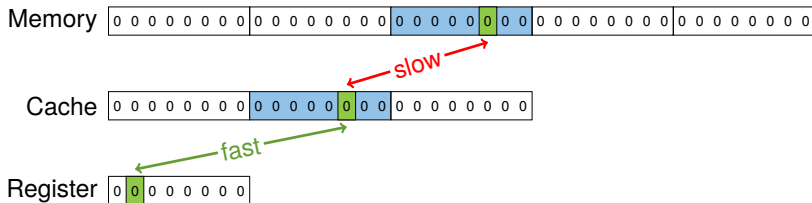
Cache organization:

- The (L1-)cache can hold multiple memory blocks
- Cache lines $\approx 100\text{kB}$



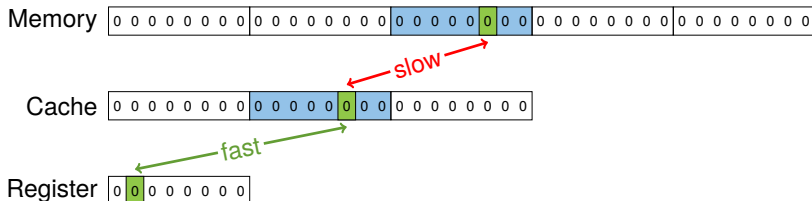
Cache organization:

- The (L1-)cache can hold multiple memory blocks
 - Cache lines $\approx 100\text{ kB}$
- If the capacity is reached unused blocks are discarded



Cache organization:

- The (L1-)cache can hold multiple memory blocks
 - Cache lines $\approx 100\text{ kB}$
- If the capacity is reached unused blocks are discarded
 - Least recently used (LRU)



Cache organization:

- The (L1-)cache can hold multiple memory blocks
 - Cache lines ≈ 100 kB
- If the capacity is reached unused blocks are discarded
 - Least recently used (LRU)
 - Least frequently used (LFU)



Cache organization:

- The (L1-)cache can hold multiple memory blocks
 - Cache lines $\approx 100\text{ kB}$
- If the capacity is reached unused blocks are discarded
 - Least recently used (LRU)
 - Least frequently used (LFU)
 - First in first out (FIFO)

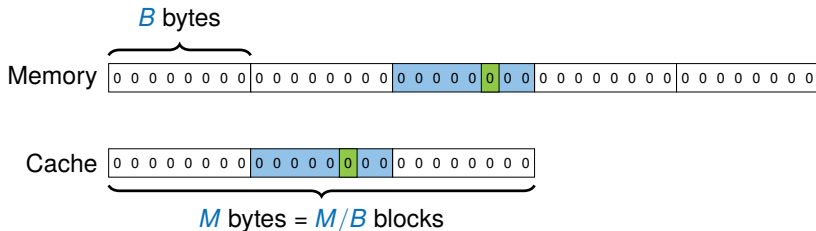


Cache organization:

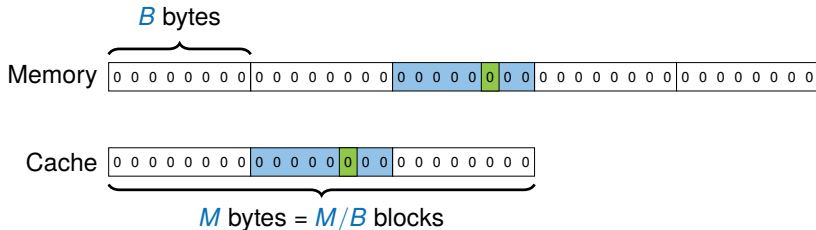
- The (L1-)cache can hold multiple memory blocks
 - Cache lines ≈ 100 kB
- If the capacity is reached unused blocks are discarded
 - Least recently used (LRU)
 - Least frequently used (LFU)
 - First in first out (FIFO)
- Details of discarding are not the topic for today

Cache Efficiency

Block Operations



Terminology:



Terminology:

- The system consists of slow and fast memory



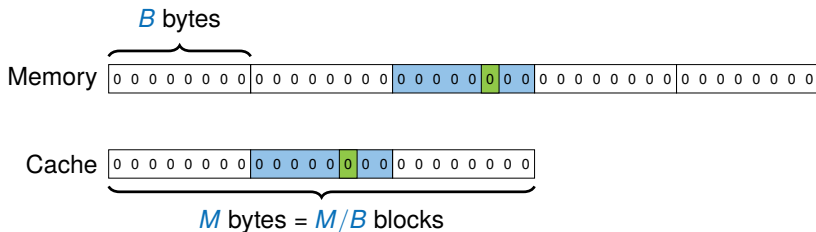
Terminology:

- The system consists of slow and fast memory
- The **slow memory** is divided in **blocks of size B**



Terminology:

- The system consists of slow and fast memory
- The **slow memory** is divided in **blocks of size B**
- The **fast cache** has size M and can store M/B blocks

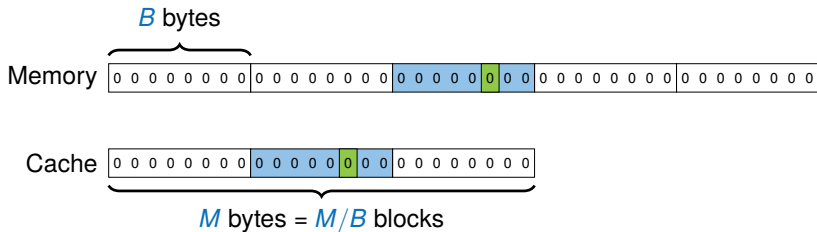


Terminology:

- The system consists of slow and fast memory
- The **slow memory** is divided in **blocks of size B**
- The **fast cache** has size M and can store M/B blocks
- If data is not in fast memory, the corresponding block is loaded into the **cache**

Cache Efficiency

Block Operations



Terminology:



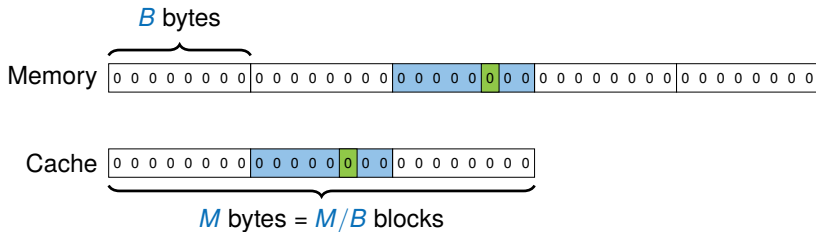
Terminology:

- The program defines which blocks are held in the **cache**



Terminology:

- The program defines which blocks are held in the **cache**
- We use the number of **block operations** as runtime estimation



Terminology:

- The program defines which blocks are held in the **cache**
- We use the number of **block operations** as runtime estimation
- We ignore runtime costs of cache accesses / management



Figure: Comparison good / bad locality

Accessing the cache B times:

- **Best case:** 1 block operation \rightarrow good locality
- **Worst case:** B block operations \rightarrow bad locality



Additional factors:

Additional factors:

- The following settings change only a small constant factor in number of block operations

Additional factors:

- The following settings change only a small constant factor in number of block operations
 - Partitioning of the slow memory into blocks

Additional factors:

- The following settings change only a small constant factor in number of block operations
 - Partitioning of the slow memory into blocks
 - Regardless of the block size: 1 Bytes or 4 Bytes or 8 Bytes

Additional factors:

- The following settings change only a small constant factor in number of block operations
 - Partitioning of the slow memory into blocks
 - Regardless of the block size: 1 Bytes or 4 Bytes or 8 Bytes

Note:

Additional factors:

- The following settings change only a small constant factor in number of block operations
 - Partitioning of the slow memory into blocks
 - Regardless of the block size: 1 Bytes or 4 Bytes or 8 Bytes

Note:

- If the input size is smaller than M we load the complete data chunk directly into the cache

Additional factors:

- The following settings change only a small constant factor in number of block operations
 - Partitioning of the slow memory into blocks
 - Regardless of the block size: 1 Bytes or 4 Bytes or 8 Bytes

Note:

- If the input size is smaller than M we load the complete data chunk directly into the cache
- Cache handling is only interesting when the input size is greater than M

Typical values: (Intel© i7-4770 Haswell, WD© Blue 2TB)

Typical values: (Intel© i7-4770 Haswell, WD© Blue 2TB)

- CPU L1 Cache: $B = 64B$, $M = 4 \times (32kB + 32kB)$

Typical values: (Intel© i7-4770 Haswell, WD© Blue 2TB)

- CPU L1 Cache: $B = 64\text{ B}$, $M = 4 \times (32\text{ kB} + 32\text{ kB})$
- CPU L2 Cache: $B = 64\text{ B}$, $M = 4 \times 256\text{ kB}$

Typical values: (Intel© i7-4770 Haswell, WD© Blue 2 TB)

- CPU L1 Cache: $B = 64\text{ B}$, $M = 4 \times (32\text{ kB} + 32\text{ kB})$
- CPU L2 Cache: $B = 64\text{ B}$, $M = 4 \times 256\text{ kB}$
- CPU L3 Cache: $B = 64\text{ B}$, $M = 8\text{ MB}$
- Disk Cache: $B = 64\text{ kB}$, $M = 64\text{ MB}$

Typical values: (Intel© i7-4770 Haswell, WD© Blue 2TB)

- CPU L1 Cache: $B = 64\text{ B}$, $M = 4 \times (32\text{ kB} + 32\text{ kB})$
- CPU L2 Cache: $B = 64\text{ B}$, $M = 4 \times 256\text{ kB}$
- CPU L3 Cache: $B = 64\text{ B}$, $M = 8\text{ MB}$
- Disk Cache: $B = 64\text{ kB}$, $M = 64\text{ MB}$
 - Many operating systems use free system memory as disk cache



Terminology:

Terminology:

- Block loads on CPU-cache are called **cache misses**

Terminology:

- Block loads on CPU-cache are called **cache misses**
- Block operations on disk-cache are called **IOs**
(input / output operations)

Terminology:

- Block loads on CPU-cache are called **cache misses**
- Block operations on disk-cache are called **IOs**
(input / output operations)
- These also fall under the term **cache efficiency** or **IO efficiency**



Example 1 - Linear order:

Example 1 - Linear order:

- We sum up all elements in **natural order**

$$\text{sum}(a) = a[1] + a[2] + \cdots + a[n]$$

Example 1 - Linear order:

- We sum up all elements in **natural order**

$$\text{sum}(a) = a[1] + a[2] + \dots + a[n]$$

- The number of block operations is $\text{ceil}\left(\frac{n}{B}\right)$

Example 1 - Linear order:

- We sum up all elements in **natural order**

$$\text{sum}(a) = a[1] + a[2] + \dots + a[n]$$

- The number of block operations is $\text{ceil}(\frac{n}{B})$

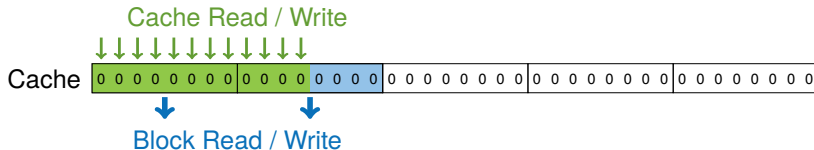


Figure: Good locality of sum operation

Example 2 - Random order:

Example 2 - Random order:

- We sum up all elements in **random order**

$$\text{sum}(a) = a[21] + a[5] + \dots + a[8]$$

Example 2 - Random order:

- We sum up all elements in **random order**

$$\text{sum}(a) = a[21] + a[5] + \dots + a[8]$$

- The number of block operations is n in the **worst case**

Example 2 - Random order:

- We sum up all elements in **random order**

$$\text{sum}(a) = a[21] + a[5] + \dots + a[8]$$

- The number of block operations is n in the **worst case**
- This leads to a runtime factor difference of B

Example 2 - Random order:

- We sum up all elements in **random order**

$$\text{sum}(a) = a[21] + a[5] + \dots + a[8]$$

- The number of block operations is n in the **worst case**
- This leads to a runtime factor difference of B



Figure: Bad locality of sum operation



Generally the factor is substantially $< B$

Generally the factor is substantially $< B$

- Even with a **random order** we access 4 neighboring bytes at once per `int` (`int32_t`)

Generally the factor is substantially $< B$

- Even with a **random order** we access 4 neighboring bytes at once per `int` (`int32_t`)
- The next element might already be loaded in the cache

Generally the factor is substantially $< B$

- Even with a **random order** we access 4 neighboring bytes at once per `int` (`int32_t`)
- The next element might already be loaded in the cache
- If **not $n \gg M$** this might occur with a high probability



QuickSort:



QuickSort:

- **Strategy:** Divide and conquer

QuickSort:

- **Strategy:** Divide and conquer
- Divide the data into two parts where the “left” part contains all values \leq those in the right part



QuickSort:

- **Strategy:** Divide and conquer
- Divide the data into two parts where the “left” part contains all values \leq those in the right part
- Choose one element (e.g the first one) as “pivot”-element



QuickSort:

- **Strategy:** Divide and conquer
- Divide the data into two parts where the “left” part contains all values \leq those in the right part
- Choose one element (e.g the first one) as “pivot”-element
- Ideally both parts are the same size

QuickSort:

- **Strategy:** Divide and conquer
- Divide the data into two parts where the “left” part contains all values \leq those in the right part
- Choose one element (e.g the first one) as “pivot”-element
- Ideally both parts are the same size
- Both parts are sorted recursively

QuickSort:

- **Strategy:** Divide and conquer
- Divide the data into two parts where the “left” part contains all values \leq those in the right part
- Choose one element (e.g the first one) as “pivot”-element
- Ideally both parts are the same size
- Both parts are sorted recursively



Figure: QuickSort with pivot-element

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



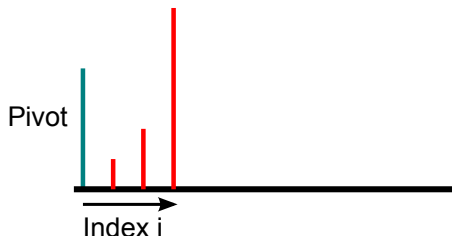
- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



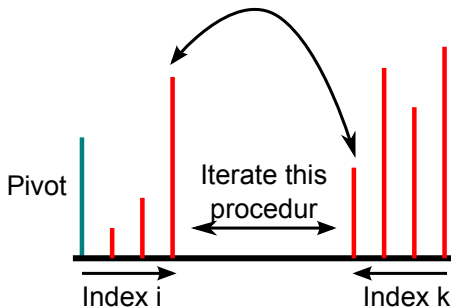
- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

- **At start:** Pivot in first position, first re-arrange list such that left part contains small, right part larger elements
- Do required changes *in place*



- **End point:** k is left to left-most element greater than pivot
swap position 0 (pivot) with k (smaller than pivot)

Python:

```
def quicksort(l, start, end):  
    if (end - start) < 1:  
        return  
  
    i = start  
    k = end  
    piv = l[0]  
  
    ...
```

```
def quicksort(l, start, end):  
    ...  
  
    while k > i:  
        while l[i] <= piv and i <= end and k > i:  
            i += 1  
        while l[k] > piv and k >= start and k >= i:  
            k -= 1  
  
        if k > i: # swap elements  
            (l[i], l[k]) = (l[k], l[i])  
  
    (l[start], l[k]) = (l[k], l[start])  
    quicksort(l, start, k - 1)  
    quicksort(l, k + 1, end)
```

Number of operations for Quicksort:

Number of operations for Quicksort:

- Let $T(n)$ be the runtime for the input size n

Number of operations for Quicksort:

- Let $T(n)$ be the runtime for the input size n

Assumptions:

Number of operations for Quicksort:

- Let $T(n)$ be the runtime for the input size n

Assumptions:

- Fields are always separated perfectly in the middle

Number of operations for Quicksort:

- Let $T(n)$ be the runtime for the input size n

Assumptions:

- Fields are always separated perfectly in the middle
- n is a power of two and recursion depth is $k = \log_2 n$

$$\begin{aligned} T(n) &\leq \underbrace{A \cdot n}_{\text{splitting in two parts}} + \underbrace{2 \cdot T\left(\frac{n}{2}\right)}_{\text{recursive sort}} \\ &\leq A \cdot n + 2 \left(A \cdot \frac{n}{2} + 2 \cdot T\left(\frac{n}{4}\right) \right) \\ &= 2A \cdot n + 4 \cdot T\left(\frac{n}{4}\right) \\ &\leq 3A \cdot n + 8 \cdot T\left(\frac{n}{8}\right) \\ &\leq \dots \\ &\leq k \cdot A \cdot n + 2^k \cdot T(1) \\ &= \log_2 n \cdot A \cdot n + n \cdot T(1) \\ &\leq \log_2 n \cdot A \cdot n + n \cdot A \in \mathcal{O}(n \log_2 n) \end{aligned}$$

Cache Efficiency

Block Operations - QuickSort

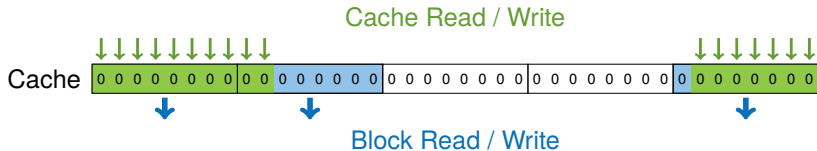


Figure: Locality of quicksort

Cache Efficiency

Block Operations - QuickSort



Figure: Locality of quicksort

- Let $IO(n)$ be the number of **block operations** for input size n



Figure: Locality of quicksort

- Let $IO(n)$ be the number of **block operations** for input size n
- Assumptions as before but recursion depth is $k = \log_2 \frac{n}{B}$
Why?

$$\begin{aligned} IO(n) &\leq \underbrace{A \cdot n/B}_{\text{splitting in two parts}} + \underbrace{2 \cdot IO(n/2)}_{\text{recursive sort}} \\ &\leq A \cdot n/B + 2(A \cdot n/2B + 2 \cdot IO(n/4)) \\ &\leq 2 \cdot A \cdot n/B + 4 \cdot IO(n/4) \\ &\leq 3 \cdot A \cdot n/B + 8 \cdot IO(n/8) \\ &\leq \dots \\ &\leq k \cdot A \cdot n/B + 2^k \cdot IO(n/2^k) \\ &= \log_2(n/B) \cdot A \cdot (n/B) + n/B \cdot IO(B) \\ &\leq \log_2(n/B) \cdot A \cdot (n/B) + A \cdot n/B \in \mathcal{O}\left(\frac{n}{B} \cdot \log_2\left(\frac{n}{B}\right)\right) \end{aligned}$$

Cache Efficiency

Introduction

Cache Organization

Divide and Conquer

Introduction

Divide and Conquer

Introduction



UNI
FREIBURG

Concept:

Concept:

- **Divide** the problem into smaller subproblems

Concept:

- **Divide** the problem into smaller subproblems
- **Conquer** the subproblems through recursive solving.
If subproblems are small enough solve them directly

Concept:

- **Divide** the problem into smaller subproblems
- **Conquer** the subproblems through recursive solving.
If subproblems are small enough solve them directly
- **Connect** all solutions of the subproblems to a solution of the full problem

Concept:

- **Divide** the problem into smaller subproblems
- **Conquer** the subproblems through recursive solving.
If subproblems are small enough solve them directly
- **Connect** all solutions of the subproblems to a solution of the full problem
- **Recursive** application of the algorithm to ever smaller subproblems

Concept:

- **Divide** the problem into smaller subproblems
- **Conquer** the subproblems through recursive solving.
If subproblems are small enough solve them directly
- **Connect** all solutions of the subproblems to a solution of the full problem
- **Recursive** application of the algorithm to ever smaller subproblems
- **Direct** solving of sufficiently small subproblems

Divide and Conquer

Introduction - Python



**UNI
FREIBURG**



- Function `solve` for solving a `problem` of size `n`

- Function `solve` for solving a problem of size n

```
def solve(problem):  
    if n < threshold:  
        return solution # solve directly  
    else:  
        # divide problem into subproblems  
        # P1, P2, ..., Pk with k>=2  
        S1 = solve(P1)  
        S2 = solve(P2)  
        ...  
        Sk = solve(Pk)  
  
        # combine solutions  
    return S1 + S2 + ... + Sk
```

Divide and Conquer

Features



UNI
FREIBURG

Divide and Conquer:



Divide and Conquer:

- Can help with conceptual hard problems

Divide and Conquer:

- Can help with conceptual hard problems
- **Solution** of the trivial problems has to be known

Divide and Conquer:

- Can help with conceptual hard problems
- **Solution** of the trivial problems has to be known
- **Dividing** in subproblems has to be possible

Divide and Conquer:

- Can help with conceptual hard problems
- **Solution** of the trivial problems has to be known
- **Dividing** in subproblems has to be possible
- **Combination** of solutions has to be possible



Features:

- Realization of **efficient solutions**



Features:

- Realization of efficient solutions
 - If trivial solution is $\in O(1)$

Features:

- Realization of efficient solutions
 - If trivial solution is $\in O(1)$
 - And separation / combination of subproblems is $\in O(n)$

Features:

- Realization of efficient solutions
 - If trivial solution is $\in O(1)$
 - And separation / combination of subproblems is $\in O(n)$
 - And the number of subproblems is limited

Features:

- Realization of **efficient solutions**
 - If trivial solution is $\in O(1)$
 - And separation / combination of subproblems is $\in O(n)$
 - And the number of subproblems is limited
 - The runtime is $\in O(n \cdot \log n)$

Features:

- Realization of **efficient solutions**
 - If trivial solution is $\in O(1)$
 - And separation / combination of subproblems is $\in O(n)$
 - And the number of subproblems is limited
 - The runtime is $\in O(n \cdot \log n)$
- Suitable for parallel processing

Features:

- Realization of **efficient solutions**
 - If trivial solution is $\in O(1)$
 - And separation / combination of subproblems is $\in O(n)$
 - And the number of subproblems is limited
 - The runtime is $\in O(n \cdot \log n)$
- Suitable for parallel processing
 - Subproblems are **independent** of each other

Features:

- Realization of **efficient solutions**
 - If trivial solution is $\in O(1)$
 - And separation / combination of subproblems is $\in O(n)$
 - And the number of subproblems is limited
 - The runtime is $\in O(n \cdot \log n)$
- Suitable for parallel processing
 - Subproblems are **independent** of each other
 - Only needed input for each subproblem has to be known



Definition of the trivial case:

Definition of the trivial case:

- Smaller subproblems are elegant and simple

Definition of the trivial case:

- Smaller subproblems are elegant and simple
- Otherwise the efficiency will be improved if relative big subproblems can be solved directly

Definition of the trivial case:

- Smaller subproblems are elegant and simple
- Otherwise the efficiency will be improved if relative big subproblems can be solved directly
- Recursion depth should not get too big (stack / memory overhead)

Division in subproblems:

Division in subproblems:

- Choosing the number of subproblems and the concrete allocation can be demanding

Division in subproblems:

- Choosing the number of subproblems and the concrete allocation can be demanding

Combination of solutions:

Division in subproblems:

- Choosing the number of subproblems and the concrete allocation can be demanding

Combination of solutions:

- Typically conceptual demanding

Divide and Conquer

Example - Maximum Subtotal



Example - Maximum Subtotal

Divide and Conquer

Example - Maximum Subtotal



Example - Maximum Subtotal Input:

Divide and Conquer

Example - Maximum Subtotal



Example - Maximum Subtotal Input:



Example - Maximum Subtotal Input:

- Sequence X of n integers



Example - Maximum Subtotal Input:

- Sequence X of n integers

Output:

Example - Maximum Subtotal Input:

- Sequence X of n integers

Output:

- Maximum sum of related subsequence and its index boundary

Example - Maximum Subtotal Input:

- Sequence X of n integers

Output:

- Maximum sum of related subsequence and its index boundary

Index	0	1	2	3	4	5	6	7	8	9
Value	31	-41	59	26	-53	58	97	-93	-23	84

Example - Maximum Subtotal Input:

- Sequence X of n integers

Output:

- Maximum sum of related subsequence and its index boundary

Index	0	1	2	3	4	5	6	7	8	9
Value	31	-41	59	26	-53	58	97	-93	-23	84

Output: Sum: 187, Start: 2, End: 6

Application:

- Maximum profit of buying and selling shares



Figure: Stock value over time

Divide and Conquer

Example - Maximum Subtotal - Python



UNI
FREIBURG

Naive solution (brute force)

Naive solution (brute force)

```
def maxSubArray(X):  
    # Store sum, start, end  
    result = (X[0], 0, 0)  
    for i in range(0, len(X)):  
        for j in range(i, len(X)):  
            subSum = 0  
            for k in range(i, j + 1):  
                subSum += X[k]  
            if result[0] < subSum:  
                result = (subSum, i, j)  
    return result
```

Divide and Conquer

Example - Maximum Subtotal - Python



UNI
FREIBURG

Runtime - Upper bound

Runtime - Upper bound

```
def maxSubArray(X):  
    result = (X[0], 0, 0)  
    # n loops -> O(n)  
    for i in range(0, len(X)):  
        # max n loops -> O(n)  
        for j in range(i, len(X)):  
            # max n loops -> O(n)  
            subSum = sum(X[i:j+1])  
            if result[0] < subSum: # O(1)  
                result = (subSum, i, j)  
    return result
```

Divide and Conquer

Example - Maximum Subtotal

Upper bound:

Divide and Conquer

Example - Maximum Subtotal

Upper bound:

- Three interleaved loops

Upper bound:

- Three interleaved loops
- Each loop with runtime $O(n)$

Upper bound:

- Three interleaved loops
- Each loop with runtime $O(n)$
- Algorithm runtime of $O(n^3)$

Lower bound:

Table: Operations

i	Additions	j
$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$

Lower bound:

Table: Operations

i	Additions	j
$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$

- We iterate at least $\frac{n}{3}$ values for i

Lower bound:

Table: Operations

i	Additions	j
$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$

- We iterate at least $\frac{n}{3}$ values for i
- For each i we iterate at least $\frac{n}{3}$ values for j

Lower bound:

Table: Operations

i	Additions	j
$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$

- We iterate at least $\frac{n}{3}$ values for i
- For each i we iterate at least $\frac{n}{3}$ values for j
- For each j we have at least $\frac{n}{3}$ additions

Lower bound:

Table: Operations

i	Additions	j
$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$	$\frac{n}{3} \in O(n)$

- We iterate at least $\frac{n}{3}$ values for i
- For each i we iterate at least $\frac{n}{3}$ values for j
- For each j we have at least $\frac{n}{3}$ additions
- We need at least $T(n) = (\frac{n}{3})^3 \in \Omega(n^3)$ steps

Divide and Conquer

Example - Maximum Subtotal - Runtime



UNI
FREIBURG

Runtime:

Runtime:

- With $T(n) \in O(n^3)$ and $T(n) \in \Omega(n^3)$ we know:

$$T(n) \in \Theta(n^3)$$

Runtime:

- With $T(n) \in O(n^3)$ and $T(n) \in \Omega(n^3)$ we know:

$$T(n) \in \Theta(n^3)$$

- It is hard to solve the problem in a worse way ...

Divide and Conquer

Example - Maximum Subtotal - Runtime



UNI
FREIBURG

Current approach:

Current approach:

- Calculating the sum for range from i to j with loop

$$S_{i,j} = X[i] + X[i+1] + \dots + X[j]$$

Current approach:

- Calculating the sum for range from i to j with loop

$$S_{i,j} = X[i] + X[i+1] + \dots + X[j]$$

Better approach:

Current approach:

- Calculating the sum for range from i to j with loop

$$S_{i,j} = X[i] + X[i+1] + \dots + X[j]$$

Better approach:

- Incremental sum instead of loop

$$S_{i,j+1} = X[i] + X[i+1] + \dots + X[j] + X[j+1]$$

$$S_{i,j+1} = S_{i,j} + X[j+1] \in O(1) \quad \text{instead of} \quad \in O(n)$$

Divide and Conquer

Example - Maximum Subtotal - Python



UNI
FREIBURG

Better solution:

Better solution:

```
def maxSubArray(X):  
    result = (X[0], 0, 0)  
    # n loops -> O(n)  
    for i in range(0, len(X)):  
        subSum = 0  
        # max n loops -> O(n)  
        for j in range(i, len(X)):  
            subSum += X[j] # O(1)  
            if result[0] < subSum: # O(1)  
                result = (subSum, i, j)  
    return result
```

Better solution:

```
def maxSubArray(X):  
    result = (X[0], 0, 0)  
    # n loops -> O(n)  
    for i in range(0, len(X)):  
        subSum = 0  
        # max n loops -> O(n)  
        for j in range(i, len(X)):  
            subSum += X[j] # O(1)  
            if result[0] < subSum: # O(1)  
                result = (subSum, i, j)  
    return result
```

■ Runtime $\in O(n^2)$

Divide and Conquer

Example - Maximum Subtotal

Divide and Conquer:



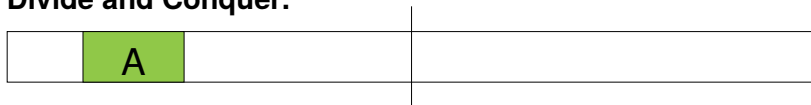
Divide and Conquer Idea to solve:

- Split the sequence in the middle

Divide and Conquer

Example - Maximum Subtotal

Divide and Conquer:



Divide and Conquer Idea to solve:

- Split the sequence in the middle
- Solve left half of the problem

Divide and Conquer:



Divide and Conquer Idea to solve:

- Split the sequence in the middle
- Solve left half of the problem
- Solve right half and combine both solutions into one

Divide and Conquer:



Divide and Conquer Idea to solve:

- Split the sequence in the middle
- Solve left half of the problem
- Solve right half and combine both solutions into one
- Maximum might be located in **left half (A)** or **right half (B)**

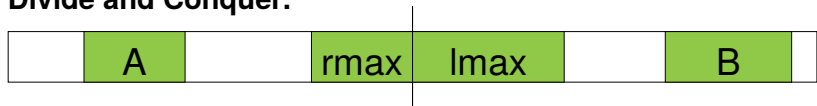
Divide and Conquer:



Divide and Conquer Idea to solve:

- Split the sequence in the middle
- Solve left half of the problem
- Solve right half and combine both solutions into one
- Maximum might be located in **left half (A)** or **right half (B)**
- Problem: Maximum can **overlap the split**

Divide and Conquer:



Divide and Conquer Idea to solve:

- Split the sequence in the middle
- Solve left half of the problem
- Solve right half and combine both solutions into one
- Maximum might be located in left half (A) or right half (B)
- Problem: Maximum can overlap the split
- To solve this case we have to calculate rmax and lmax

Divide and Conquer:



Divide and Conquer Idea to solve:

- Split the sequence in the middle
- Solve left half of the problem
- Solve right half and combine both solutions into one
- Maximum might be located in **left half (A)** or **right half (B)**
- Problem: Maximum can **overlap the split**
- To solve this case we have to calculate **rmax** and **lmax**
- The overall solution is the **maximum of A, B and C**

Divide and Conquer

Example - Maximum Subtotal



Principle - Divide and Conquer:

Principle - Divide and Conquer:

- Small problems are solved directly: $n = 1 \Rightarrow \text{max} = X[0]$

Principle - Divide and Conquer:

- Small problems are solved directly: $n = 1 \Rightarrow \text{max} = X[0]$
- Bigger problems are partitioned into two subproblems and recursively solved. Subsolutions **A** and **B** are returned

Principle - Divide and Conquer:

- Small problems are solved directly: $n = 1 \Rightarrow \text{max} = X[0]$
- Bigger problems are partitioned into two subproblems and recursively solved. Subsolutions **A** and **B** are returned
- To determine subsolution **C**, **rmax** and **lmax** for the subproblems are computed

Principle - Divide and Conquer:

- Small problems are solved directly: $n = 1 \Rightarrow \text{max} = X[0]$
- Bigger problems are partitioned into two subproblems and recursively solved. Subsolutions **A** and **B** are returned
- To determine subsolution **C**, **rmax** and **lmax** for the subproblems are computed
- The overall solution is the **maximum of A, B and C**

```
def maxSubArray(X, i, j):  
    if i == j: # trivial case  
        return (X[i], i, i)  
  
    # recursive subsolutions for A, B  
    m = (i + j) / 2  
    A = maxSubArray(X, i, m)  
    B = maxSubArray(X, m + 1, j)  
  
    # rmax and lmax for cornercase C  
    C1, C2 = rmax(X, i, m), lmax(X, m + 1, j)  
    C = (C1[0] + C2[0], C1[1], C2[1])  
  
    # compute solution from results A, B, C  
    return max([A, B, C], key=lambda i: i[0])
```

■ General

[CRL01] Thomas H. Cormen, Ronald L. Rivest, and Charles E. Leiserson.

Introduction to Algorithms.

MIT Press, Cambridge, Mass, 2001.

[MS08] Kurt Mehlhorn and Peter Sanders.

Algorithms and data structures, 2008.

<https://people.mpi-inf.mpg.de/~mehlhorn/ftp/Mehlhorn-Sanders-Toolbox.pdf>.

■ Caching

[Wik] [Cache](https://en.wikipedia.org/wiki/Cache)

`https://en.wikipedia.org/wiki/Cache`