Algorithms and Datastructures
Balanced Trees (AVL-Trees, (a,b)-Trees, Red-Black-Trees)

Albert-Ludwigs-Universität Freiburg

Prof. Dr. Rolf Backofen

Bioinformatics Group / Department of Computer Science
Algorithms and Datastructures, January 2017

UNI
FREIBURG

Balanced Trees
Motivation

Binary search tree:

- With `BinarySearchTree` we could perform an `lookup` or `insert` in $O(d)$, with d being the depth of the tree
- Best case: $d \in O(\log n)$, keys are inserted randomly
- Worst case: $d \in O(n)$, keys are inserted in ascending / descending order $(20, 19, 18, \ldots)$

Gnarley trees:

- http://people.ksp.sk/~kuko/bak

Figure: Binary search tree with random insert [Gna]

Figure: Binary search tree with descending insert [Gna]

Balanced trees:

- We do not want to rely on certain properties of our key set
- We explicitly want a depth of $O(\log n)$
- We rebalance the tree from time to time

Balanced Trees
Motivation

How do we get a depth of $O(\log n)$?

- **AVL-Tree:**
  - Binary tree with 2 children per node
  - Balancing via "rotation"
- **(a,b)-Tree** or **B-Tree:**
  - Node has between *a* and *b* children
  - Balancing through splitting and merging nodes
  - Used in databases and file systems
- **Red-Black-Tree:**
  - Binary tree with "black" and "red" nodes
  - Balancing through "rotation" and "recoloring"
  - Can be interpreted as (2, 4)-tree
  - Used in C++ `std::map` and Java `SortedMap`

Balanced Trees
AVL-Tree

AVL-Tree:

- Gregory Maximovich **A**delson-**V**elskii, Yevgeniy Mikhailovlovich **L**andis (1963)
- Search tree with modified `insert` and `remove` operations while satisfying a depth condition
- Prevents degeneration of the search tree
- Height difference of left and right subtree is at maximum one
- With that the height of the search tree is always $O(\log n)$
- We can perform all basic operations in $O(\log n)$

Figure: Example of an AVL-Tree

Figure: **Not** an AVL-Tree

Figure: Another example of an AVL-Tree
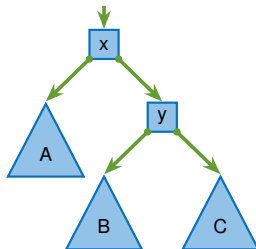
Rotation:

Figure: Before rotating

$\Rightarrow$

Figure: After rotating

- Central operation of rebalancing
- After rotation to the right:
    - Subtree *A* is a layer higher and subtree *C* a layer lower
    - The parent child relations between nodes *x* and *y* have been swapped

AVL-Tree:

- If a height difference of ± 2 occurs on an `insert` or `remove` operation the tree is rebalanced
- Many different cases of rebalancing
- **Example:** `insert` of $1, 2, 3, \ldots$

Figure: Inserting $1, \ldots, 10$ into an AVL-tree [Gna]

Summary:

- Historical the first search tree providing guaranteed `insert`, `remove` and `lookup` in $O(\log n)$
- However not amortized update costs of $O(1)$
- Additional memory costs: We have to save a height difference for every node
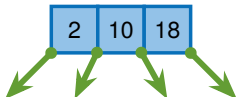- Better (and easier) to implement are (a,b)-trees

(a,b)-Tree:

- Also known as **b-tree** (b for "balanced")
- Used in databases and file systems

Idea:

- Save a varying number of elements per node
- So we have space for elements on an `insert` and balance operation

(a,b)-Trees
Introduction

(a,b)-Tree:

- All leaves have the same depth
- Each inner node has $\geq a$ and $\leq b$ nodes
  (Only the root node may have less nodes)

- Each node with n children is called "node of degree n" and holds $n-1$ sorted elements
- Subtrees are located "between" the elements
- We require: $a \geq 2$ and $b \geq 2a - 1$
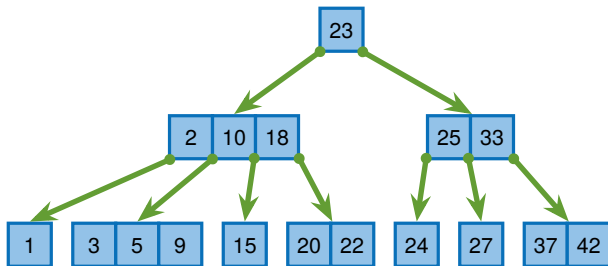
(2,4)-Tree:

Figure: Example of an (2,4)-tree

- (2,4)-tree with depth of 3
- Each node has between 2 and 4 children (1 to 3 elements)

Not an (2,4)-Tree:

Figure: **Not** an (2,4)-tree

- Invalid sorting
- Degree of node too large / too small
- Leaves on different levels

Searching an element: (`lookup`)

- The same algorithm as in BinarySearchTree
- Searching from the root downwards
- The keys at each node set the path

Figure: (3,5)-Tree [Gna]

Inserting an element: (insert)

- Search the position to insert the key into
- This position will always be an leaf
- Insert the element into the tree
- **Attention:** As a result node can overflow by one element (Degree $b + 1$)
- Then we **split** the node

Inserting an element: (insert)

Figure: Splitting a node

- If the degree is higher than $b+1$ we split the node
- This results in a node with $\mathrm{ceil}\left(\frac{b-1}{2}\right)$ elements, a node with $\mathrm{floor}\left(\frac{b-1}{2}\right)$ elements and one element for the parent node
- Thats why we have the limit $b \geq 2a - 1$

Inserting an element: (`insert`)

- If the degree is higher than $b+1$ we split the node
- Now the parent node can be of a higher degree than $b+1$
- We split the parent nodes the same way
- If we split the root node we create a new parent root node
  (The tree is now one level deeper)

Removing an element: (`remove`)

- Search the element in $O(\log n)$ time
- **Case 1:** The element is contained by a leaf
  - Remove element
- **Case 2:** The element is contained by an inner node
  - Search the successor in the right subtree
  - The successor is always contained by a leaf
  - Replace the element with its successor and delete the successor from the leaf
- **Attention:** The leaf might be too small (degree of $a - 1$)
  $\Rightarrow$ We rebalance the tree

Removing an element: (`remove`)

- **Attention:** The leaf might be too small (degree of $a - 1$)
  $\Rightarrow$ We rebalance the tree

  - **Case a:** If the left or right neighbour node has a degree greater than a we **borrow** one element from this node

Figure: Borrow an element

Removing an element: (`remove`)

- **Attention:** The leaf might be too small (degree of $a - 1$)
  $\Rightarrow$ We rebalance the tree

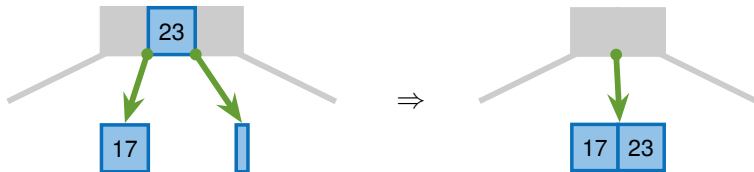 - **Case b:** We merge the node with its right or left neighbour

Figure: Merge two nodes

Removing an element: (`remove`)

- Now the parent node can be of degree $a - 1$
- We merge parent nodes the same way
- If the root has only a single child
    - Remove the root
    - Define sole child as new root
    - The tree shrinks by one level

Runtime complexity of `lookup`, `insert` and `remove`:

- All operations in $O(d)$ with d being the depth of the tree
- Each node (except the root) has more than a children
  $\Rightarrow n \geq a^{d-1}$ and $d \leq 1 + \log_a n = O(\log_a n)$

In detail:

- `lookup` always takes $\Theta(d)$
- `insert` and `remove` often require only $O(1)$ time
- **Worst case:** split or merge all nodes on path up to the root
- Therefore instead of $b \geq 2a - 1$ we need $b \geq 2a$

**Counter example (2,3)-Tree:**

- Before executing `delete(11)`

Figure: Normal (2,3)-Tree

Counter example (2,3)-Tree:

- Executing delete(11)

Figure: (2,3)-Tree - Delete step 1

Counter example (2,3)-Tree:

- Executing `delete(11)`

Figure: (2,3)-Tree - Delete step 2

Counter example (2,3)-Tree:

- Executing `delete(11)`

Figure: (2,3)-Tree - Delete step 3

Counter example (2,3)-Tree:

- Executed `delete(11)`

Figure: (2,3)-Tree - Delete step 4

**Counter example (2,3)-Tree:**

- Executing `insert(11)`

Figure: (2,3)-Tree - Insert step 1

**Counter example (2,3)-Tree:**

- Executing `insert(11)`

Figure: (2,3)-Tree - Insert step 2

Counter example (2,3)-Tree:

- Executing `insert(11)`

Figure: (2,3)-Tree - Insert step 3

Counter example (2,3)-Tree:

- Executed `insert(11)`

Figure: (2,3)-Tree - Insert step 4

Counter example (2,3)-Tree:

- We are exactly where we started
- If $b = 2a - 1$ then we can create a sequence of `insert` and `remove` operations where each operation costs $O(\log n)$
- We need $b \geq 2a$ instead of $b \geq 2a - 1$

Figure: (2,3)-Tree

(2,4)-Tree:

- If all nodes have 2 children we have to merge the nodes up to the root on a `remove` operation
- If all nodes have 4 children we have to split the nodes up to the root on a `insert` operation
- If all nodes have 3 children it takes some time to reach one of the previous two states
- ⇒ **Nodes of degree 3 are stable**
  Neither an `insert` nor a `remove` operation trigger rebalancing operations

(2,4)-Tree:

- **Idea:**
    - After an expensive operation the tree is in a stable state
    - It takes some time until the next expensive operation occurs
- Like with dynamic arrays:
    - Reallocation is expensive but it takes some time until the next expensive operation occurs
    - If we overallocate clever we have an amortized runtime of $O(1)$

Terminology:

- We analyze a sequence of n operations
- Let Φ_i be the potential of the tree after the *i-th* operation
- Φ_i = the number of stable nodes with degree 3
- Empty tree has 0 nodes: $\Phi = 0$

Example:

■ Nodes of degree 3 are highlighted

Figure: Tree with potential Φ = 4

(a,b)-Trees
Runtime Complexity - (2,4)-Tree

Terminology:

- Let c_i be the costs = runtime of the i-th operation
- We will show:
  - Each operation can at most destroy one stable node
  - For each cost incurring step the operation creates an additional stable node
- The costs for operation i are coupled to the difference of the potential levels

$$c_i \leq A \cdot \underbrace{(\Phi_i - \Phi_{i-1})}_{} + B, \quad A > 0 \text{ and } B > A$$

  Number of gained stable nodes (degree 3) ≥ -1

- Each operation has an amortitzed cost of $O(1)$ summing up to $O(n)$ in total

Case 1: *i-th* operation is an `insert` operation on a full node

Figure: Splitting a node on `insert`

- Each splitted node creates a node of degree 3
- The parent node receives an element from the splitted node
- If the parent node is also full we have to split it too

Case 1: *i-th* operation is an `insert` operation on a full node

- Let m be the number of nodes split
- The potential rises by m
- If the "stop-node" is of degree 3 then the potential goes down by one

$$\Phi_i \geq \Phi_{i-1} + m - 1$$
$$\Rightarrow m \leq \Phi_i - \Phi_{i-1} + 1$$

Costs: $c_i \leq A \cdot m + B$

$$\Rightarrow c_i \leq A \cdot (\Phi_i - \Phi_{i-1} + 1) + B$$
$$c_i \leq A \cdot (\Phi_i - \Phi_{i-1}) + \underbrace{A + B}_{B'}$$

(a,b)-Trees
Runtime Complexity - (2,4)-Tree

Case 2: *i-th* operation is an `remove` operation

- **Case 2.1:** Inner node
    - Searching the successor in a tree is $O(d) = O(\log n)$
    - Normally the tree is coupled with a doubly linked list
      $\Rightarrow$ We can find the succcessor in $O(1)$

Figure: Tree with doubly linked list

Case 2: *i-th* operation is an `remove` operation

- **Case 2.1:** Borrow a node
  - Creates no additional operations
  - Case 2.1.1: Potential rises by one

Figure: Case 2.1.1: Borrow an element

Case 2: *i-th* operation is an `remove` operation

- **Case 2.1:** Borrow a node
  - Creates no additional operations
  - Case 2.1.2: Potential is lowered by one

Figure: Case 2.1.2: Borrow an element

Case 2: *i-th* operation is an `remove` operation

- **Case 2.2:** Merging two node

Figure: Merging two nodes

- Potential rises by one
- Parent node has one element less after the operation
- This operation propagates upwards until a node of degree > 2 or a node of degree 2, which can borrow from a neighbour

Case 2: *i-th* operation is an `remove` operation

- **Case 2.2:** Merging two node

Figure: Merging two nodes

- The potential rises by *m*
- If the "stop-node" is of degree 2 then the potential eventually goes down by one
- Same costs as `insert`

Lemma:

- We know:

$$c_i \leq A \cdot (\Phi_i - \Phi_{i-1}) + B, \quad A > 0 \text{ and } B > A$$

- With that we can conclude:

$$\sum_{i=0}^{n} c_i \in O(n)$$

Proof:

$$\sum_{i=0}^{n} c_i \leq \underbrace{A \cdot (\Phi_1 - \Phi_0) + B}_{\leq c_1} + \underbrace{A \cdot (\Phi_2 - \Phi_1) + B}_{\leq c_2} + \cdots + \underbrace{A \cdot (\Phi_n - \Phi_{n-1}) + B}_{\leq c_n}$$

$$= A \cdot (\Phi_n - \Phi_0) + B \cdot n \quad | \text{ telescope sum}$$

$$= A \cdot \Phi_n + B \cdot n \quad | \text{ we start with an empty tree}$$

$$< A \cdot n + B \cdot n \in O(n) \quad | \text{ number of degree 3 nodes}$$

$$< \text{ number of nodes}$$

Red-Black Tree:

- Binary tree with red and black nodes
- Number of black nodes on path to leaves is equal
- Can be interpreted as (2,4)-tree (also named 2-3-4-tree)
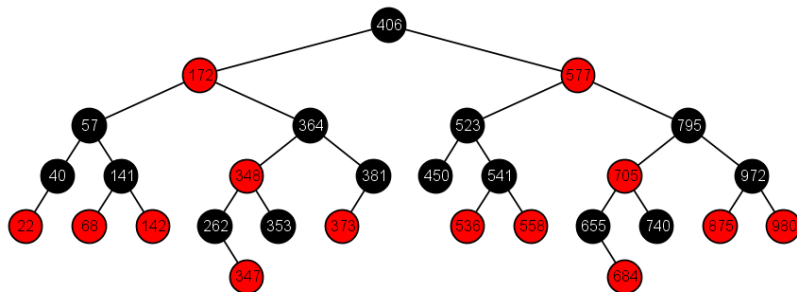- Each (2,4)-tree-node is a small red-black-tree with a black root node

Figure: Example of an red-black-tree [Gna]

- **General**

  [CRL01] Thomas H. Cormen, Ronald L. Rivest, and
          Charles E. Leiserson.
          Introduction to Algorithms.
          MIT Press, Cambridge, Mass, 2001.

  [MS08] Kurt Mehlhorn and Peter Sanders.
          Algorithms and data structures, 2008.
          `https://people.mpi-inf.mpg.de/~mehlhorn/`
          `ftp/Mehlhorn-Sanders-Toolbox.pdf`.

- **Gnarley Trees**

  [Gna]   Gnarley Trees
          `https://people.ksp.sk/~kuko/gnarley-trees/`

- **AVL-Tree**

  [Wik] AVL tree

  https://en.wikipedia.org/wiki/AVL_tree

- **(a,b)-Tree**

  [Wika] 2-3-4 tree

  https://en.wikipedia.org/wiki/2%E2%80%933%E2%80%934_tree

  [Wikb] (a,b)-tree

  https://en.wikipedia.org/wiki/(a,b)-tree

- **Red-Black-Tree**

  [Wik] Red-black tree
      https://en.wikipedia.org/wiki/Red%E2%80%93black_tree