

# Algorithms and Datastructures

## Shortest Path, Dijkstra Algorithm

Albert-Ludwigs-Universität Freiburg



**UNI  
FREIBURG**

Prof. Dr. Rolf Backofen

Bioinformatics Group / Department of Computer Science  
Algorithms and Datastructures, March 2016

Feedback

Exercises

Lecture

Graphs

Dijkstra Algorithm

# Feedback from the exercises



# Feedback from the lecture

For a graph  $G = (V, E)$ :

- A path of  $G$  is a sequence of edges  $u_1, u_2, \dots, u_i \in V$  with
  - Undirected graph:  $\{u_1, u_2\}, \{u_2, u_3\}, \dots, \{u_{i-1}, u_i\} \in E$
  - Directed graph:  $(u_1, u_2), (u_2, u_3), \dots, (u_{i-1}, u_i) \in E$
- The length of a path is
  - Without weights: number of edges taken
  - With weights: sum of weights of edges taken

For a graph  $G = (V, E)$ :

- The **shortest path** between two vertices  $u, v$  is the path  $P = (u, \dots, v)$  with the shortest length  $d(u, v)$  or lowest costs
- The **diameter** of a graph is the **longest shortest path**

# Dijkstra Algorithm



## Shortest Path without Computer

- Wanted: Shortest path from M to all other points
- Place pearls on crossings and clamp strings between them



# Dijkstra Algorithm

## Shortest Path without Computer



**Figure:** Based on OpenStreetMaps; CC BY-SA 2.0

- Take the net and pull it slowly upwards until fully lifted
- Each node (pearl) now has a specific height
- The distance to M is exactly the **shortest path**






Figure: Shortest path from  $s$  to  $t$

- Let  $r$  be the shortest path from  $s$  to  $t$
- For each node  $u$  on path  $r$  the path from  $u$  to  $t$  is the shortest path

### Proof:

- If there was a shorter path from  $s$  to  $u$  then we could choose this path to get faster to  $t$
- Then  $r$  would not be the shortest path



Figure: Shortest path from  $s$  to  $t$

- This is also correct for all sub paths on  $r$
- If the shortest path from  $s$  to  $t$  passes  $u_1$  and  $u_2$  then the sub path  $(u_1, u_2)$  is the shortest path from  $u_1$  to  $u_2$

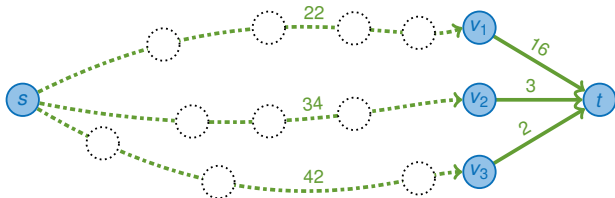


Figure: Shortest paths from  $s$  to  $t$

- If we know the shortest path from  $s$  to the preceding nodes of  $t$  ( $v_1, v_2, v_3$ ) we can determine the shortest path to  $t$

### Idea:

- Attach the cost of the shortest path to each node
- Let the information travel over the edges (message passing)
- In which order should we process the nodes?

## Inventor:

- Edsger Dijkstra (1930 - 2002)
- Computer scientist from Netherlands
- Won Turing-Award as one of few Europeans for his studies of structured programming
- Invented the Dijkstra-Algorithm in 1959



**Figure:** Portrait © Hamilton Richards - manuscripts of Edsger W. Dijkstra, University Texas at Austin

## Example:

- Lift pearl **M** a little bit
- Connections to pearls **R**, **L** and **G** are hanging in the air
- Lift further until pearl **R** starts to lift at 5m
- The shortest path to **R** is now known
- Lift further: The wires from **R**, **O** and **Q** are now in the air
- One of the pearls **G**, **L**, **Q** or **O** is the next one  
Which one?

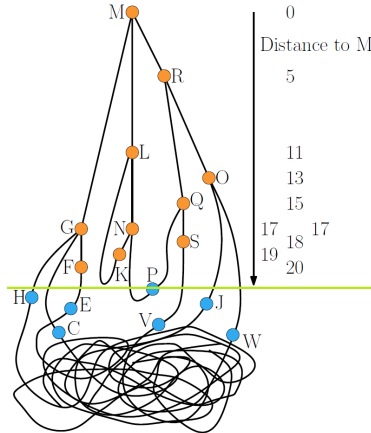


Figure: Map © Mehlhorn / Sanders

## Example:

- At 11 m pearl **L** gets lifted
- The wires to **N** and **K** are now in the air
- One of the pearls **G**, **K**, **N**, **Q** or **O** is the next one  
Which one?
- At 13 m pearl **O** gets lifted  
...
- How to translate this into an computer algorithm?







Figure: Map © Mehlhorn / Sanders

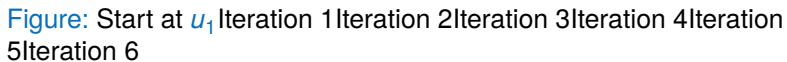
## High level description: Three types of nodes

- **Settled:** For node  $u$  we know  $\text{dist}(s, u)$   
(Pearl example: This pearl is hanging in the air) 
- **Active:** For node  $u$  we know a tentative distance  $\text{td}(u) \geq \text{dist}(s, u)$  (Can be optimal but doesn't have to)  
(Pearl example: This pearl is laying on the table but one connected wire is already in the air) 
- **Unreached:** We have not reached the node yet  
(Pearl example: This preal is hanging in the air) 



## High level description:

- Each iteration take the **active** node  $u$  with the **smallest**  $td(u)$   
(The pearl getting lifted next)
- We update the state of the node  $u$  to **settled**  
(The pearl gets lifted)
- We check for each **neighbor**  $v$  of node  $u$  if we can reach  $v$  faster than currently possible  
(Check all outgoing wires from this pearl: Activate all connected pearls, update tentative distance if smaller)
- Iterate until no active nodes exist anymore



### Proof:

- **Assumption 1:** All edges have a positive length
- **Assumption 2:** Each node has a unique distance  $\text{dist}(s, u)$   
(This was not the case on the previous slides)

This results in an easy and intuitive proof.

It is possible to show this without assumption 2. See references if interested

- With assumption 2 there exists a sorting  $u_1, u_2, \dots$  with that:

$$\text{dist}(s, u_1) < \text{dist}(s, u_2) < \text{dist}(s, u_3) < \dots$$

### Proof:

- With **assumption 2** there exists a sorting  $u_1, u_2, \dots$  with that:

$$\text{dist}(s, u_1) < \text{dist}(s, u_2) < \text{dist}(s, u_3) < \dots$$

- We want to show that the *Dijkstra* algorithm finds the shortest path for each node  $u_i$  so that  $\text{td}(u_i) = \text{dist}(s, u_i)$  holds
- Additionally we show that each node gets solved in order of the distance: Node  $u_i$  gets solved in iteration  $i$

$$u_1, u_2, u_3, \dots$$

**To show:** Node  $u_i$  gets solved in round  $i$

- 1 Node  $u_i$  contains the correct distance ( $td(u_i) = dist(s, u_i)$ ) and is active
- 2 Node  $u_i$  has the smallest value for  $td(u_i)$  and gets selected by the algorithm

**Induction start:**

- 1
  - Only the start node  $s = u_1$  is active and  $td(s) = 0$
  - Node  $u_1$  gets solved and  $td(u_1) = dist(s, u_1) = 0$
- 2 Only the start node  $u_1$  is active

### Induction step: $i = i + 1$

- 1 **To show:** Node  $u_i$  contains the correct distance ( $\text{td}(u_i) = \text{dist}(s, u_i)$ ) and is active

- On the shortest path from  $s$  to  $u_{i+1}$  is a preceding node that:

$$\text{dist}(s, u_{i+1}) = \text{dist}(s, v) + c(v, u_{i+1})$$

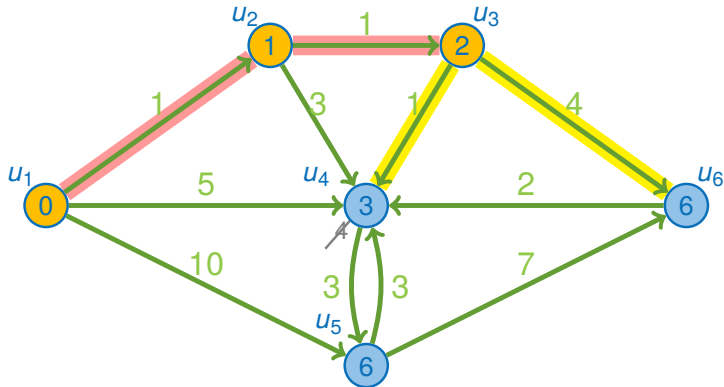
( $c$  are the costs of the edge)



- With that results  $\text{dist}(s, v) < \text{dist}(s, u_{i+1})$  because  $c > 0$
- Because  $u_{i+1}$  is currently solved node  $v$  is one of the preceding nodes  $u_1, \dots, u_i$ , hence  $v = u_j$  with  $0 \leq j \leq i$

# Dijkstra Algorithm

## Proof - Example of Iteration 6



- Preceding node of  $u_6$  is  $v = u_3$
- In round 3  $td(u_6) = 2 + 4 = 6$  was already solved



- 1 **To show:** Node  $u_i$  contains the correct distance ( $\text{td}(u_i) = \text{dist}(s, u_i)$ ) and is active
- With **induction assumption:**  $v$  already contains the correct distance which was evaluated in round  $j$  (edge from  $v$  to  $u_{i+1}$ ) and is stored in  $\text{td}(u_{i+1})$
  - $u_{i+1}$  is active because the preceding node was solved





2 **To show:** Node  $u_{i+1}$  has the smallest value for  $\text{td}(u_{i+1})$  and gets selected by the algorithm

- All nodes with smaller  $\text{dist}$  are already solved
- All other nodes  $u_k$  with  $k > i + 1$  have a greater  $\text{dist}(s, u_k)$  and with that the  $\text{td}(u_k)$  is greater or equal

$\Rightarrow u_{i+1}$  is the node with the smallest  $\text{td}$  and gets selected by the algorithm

### Implementation:

- We have to manage a set of **active nodes**
- We start with only the **start node** in our set
- At the start of each iteration we need the node  $u$  with the smallest  $td(u)$

How to implement this?

### Implementation:

- Using a **priority queue** with  $td(u)$  as keys
- The following problem occurs:
  - The **tentative distance** of an active node might change multiple times before it is settled
  - We have to change the key in our **priority queue** without removing the entry

### Limitations:

- Often only insert, getMin and deleteMin are implemented
- ⇒ We only have access to the first element and not any desired one

### Alternative:

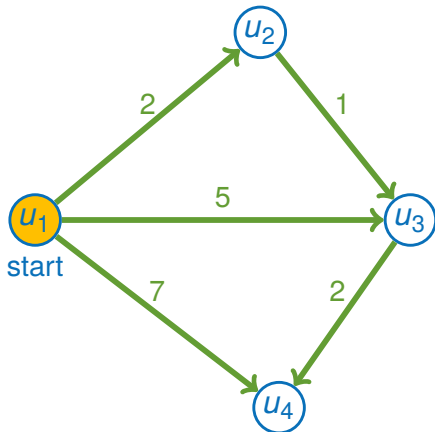
- If a node reoccurs with a smaller **dist** we insert the element one more time into the **priority queue**  
(We do nothing if the distance is greater or equal)
- We do not remove the old entry
- The node always gets solved with the smallest distance because of the priority
- If a node reoccurs with a higher **dist** we remove it and do simply **nothing**

# Dijkstra Algorithm

## Implementation - Example



Priority queue:



Graph with  $n$  nodes and  $m$  edges: ( $m \geq n$ )

- Each node gets solved exactly **one time**
- When solving a node its outgoing edges are taken into account
- Each edge triggers at maximum one `insert` operation
- The number of operations on the **priority queue** is at maximum  $O(m)$
- This results in a runtime of  $O(m \cdot \log m)$   
( $\log m$  because of at max.  $m$  elements in the priority queue)

Runtime of  $O(m \cdot \log m)$ :

- Because of  $m \leq n^2$  we have a maximum runtime of  $O(m \cdot \log n)$ , because  $\log n^2 = 2 \log n$
  - With a complex **priority queue** the runtime can be reduced to  $O(m + n \log n)$ 
    - For example with a **Fibonacci heap**
    - This results in a better runtime for complex graphs  $m \sim n^2$
    - Complex heaps create a management overhead
- ⇒ In practice  $m \in O(n)$  with a **binary heap** being faster  
(See lecture 6)

### Termination criteria:

- Terminate as soon as the target node  $t$  is settled  
... never before because tentative distance might change:

$$td(t) \geq dist(s, t)$$

- Before the node  $t$  is solved **all nodes  $u$**  with  $dist(s, u) \leq dist(s, t)$  are settled



### Termination criteria:

- Not only the **single source single target** shortest path problem is solved by the Dijkstra algorithm but also the **single source all targets** problem
- This sounds wasteful but there is not a (much) better method for general graphs **Intuitive:** We only know that there is no shorter path if all in the range of  $\text{dist}(s, t)$  around  $s$  is evaluated

### Calculate the shortest path:

- With the current implementation of the Dijkstra algorithm we only get the **length** of the path  
How to get the path too?
- If we save the preceding node of the current shortest path on **relaxation** of each node we can reconstruct the **path**

# Dijkstra Algorithm




Figure: Start at  $u_1$  Iteration 1 Iteration 2 Iteration 3 Iteration 4 Iteration 5 Iteration 6

### Enhancement:

- In our proof we used the assumption that all costs are **not negative** (even  $> 0$ )
- With **negative costs** there might be **negative cycles**:




Figure: Here no problem ...


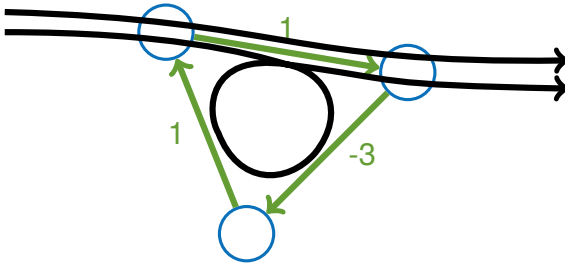


Figure: ... but here

### Negative cycles:



- No cycle:  
cost of 1
- 1 cycle:  
cost of 0
- 2 cycles:  
cost of -1
- 3 cycles:  
cost of -2
- ...

### Enhancement:

- We need a different algorithm to deal with negative edges
  - For example the **Bellman-Ford** algorithm
  - If the graph is **acyclic** we can simply use a topological sorting (with DFS) and relaxing the nodes in order of this sorting
- Another (not only) in artificial intelligence used variant of the Dijkstra algorithm is the **A\* algorithm**

Additional information given:

$h(u)$  = estimated value for  $\text{dist}(u, t)$

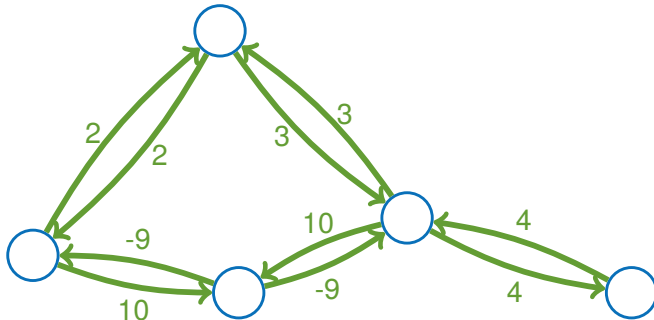
# Dijkstra Algorithm

Example - Negative costs (e-car consumption)



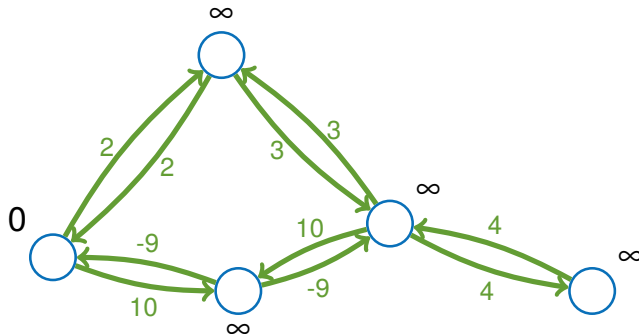
## Dijkstra algorithm:

Message passing only from solved nodes



### Bellman-Ford algorithm:

Message passing from all nodes until the path lengths are stable






### Application example:

- Route planner for car trips (exercise sheet)
- Route planner for bus / train connections

What could the graph look like?

### Space-time graph:

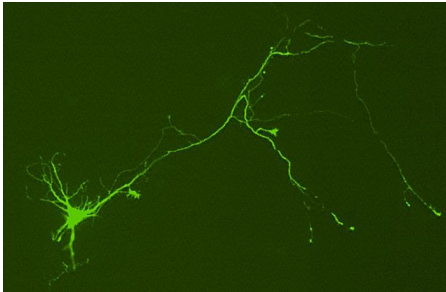


# Dijkstra Algorithm

Application in image processing



UNI  
FREIBURG




**Figure:** Neurons under fluorescence microscope

- **Task:** Measure length of axons (connections of neurons)
- Demo with ImageJ plugin NeuronJ  
<http://www.imagescience.org/meijering/software/neuronj/>

# Dijkstra Algorithm

Application: Trace axons



- Image as graph: Each pixel is a node
- Implicit edges: Each pixel has an edge to its 8 neighbours (no need to save the edges)
- Costs for nodes (not edges): bright pixels are cheap, dark pixels are costly

## ■ General

[CRL01] Thomas H. Cormen, Ronald L. Rivest, and Charles E. Leiserson.

**Introduction to Algorithms.**

MIT Press, Cambridge, Mass, 2001.

[MS08] Kurt Mehlhorn and Peter Sanders.

Algorithms and data structures, 2008.

<https://people.mpi-inf.mpg.de/~mehlhorn/ftp/Mehlhorn-Sanders-Toolbox.pdf>.

## ■ Dijkstra's algorithm

[Wik] [Dijkstra's algorithm](#)

`https:`

`//en.wikipedia.org/wiki/Dijkstra's_algorithm`

## ■ Shortest path problem

[Wik] [Shortest path problem](#)

`https://en.wikipedia.org/wiki/Shortest_path_  
problem`